

# LAB 01 - LINUX BACKUP

---

Jeremy Zerbib - Adrien Barth

---

## Task 1 : Local Sync

---

1. The output of

```
du -sh /home
210M    /home
```

The command used to create a full uncompressed backup is :

```
sudo tar cf /mnt/backup2/uncompressed.tar /home/
```

2. In order to create the backup directory, we used the following command :

```
sudo mkdir /mnt/backup2/jeremy_backup
```

Also, we changed the rights using this command :

```
sudo chown jeremy /mnt/backup2/jeremy_backup/
```

3. After typing the command :

```
rsync -av /home/ 2017-09-25-093533
```

- **What do these options do?**

Those options ensure that :

- `-v` : The console output is verbose. It explains all the operations the command is doing
- `-a` : Turns in the *archive mode*.

- **Specifically, which options are implied by the `-a` option and what do they do?**

All options are implied with `-a` with the exception of *hard links* being preserved. As this kind of operation is expensive, you want to ensure the `-H` option is specified to preserve those links.

- **How can you use the `date` command to avoid typing the timestamp of the current time? How do you make `date` produce UTC time?**

Using the `date` command, you can produce the UTC time using the `-u` option.

- **How much disk space is used by the backup directory?**

Using the `du -sh 2017-09-25-093533/` command, we find out the total space is 44MB large.

4. The command used is :

```
rsync -av /home/ 2017-09-25-093533 --delete --link-dest=2017-09-25-093533/
```

- `--delete` : Deletes files in the destination directory if they don't exist in the source directory.
- `--link-dest=DIR` : hard link to files in *DIR* when unchanged.
- **How much disk space is used by the backup directory according to the `du` command?**  
44MB are used according to the `du` command.
- **How much by the individual snapshot directories?**  
According to the individual snapshot directory, 946.4 kB are used.
- **How do you explain what `du` displays (if you had to write the `du` command, how would you count hard links)?**

We could explain that `du` reads the links and therefore, it reads the size of those links.

5. In the `jeremy` folder, we created a file called `a` using this :

```
touch a
```

We ran the `stat` command on a screenshot done some time prior to the backup

```
stat /home/jeremy/Pictures/Screenshot\ from\ 2019-09-25\ 22-01-01.png
  File: /home/jeremy/Pictures/Screenshot from 2019-09-25 22-01-01.png
  Size: 445774      Blocks: 872      IO Block: 4096   regular file
Device: 801h/2049d Inode: 131311     Links: 1
Access: (0664/-rw-rw-r--)  Uid: ( 1000/  jeremy)   Gid: ( 1000/  jeremy)
Access: 2019-10-13 13:36:16.777078803 +0200
Modify: 2019-09-25 22:01:01.653816263 +0200
Change: 2019-09-25 22:01:01.653816263 +0200
Birth: -
```

Then we ran it on the file created :

```
stat /home/jeremy/a
  File: /home/jeremy/a
  Size: 0      Blocks: 0      IO Block: 4096   regular empty file
Device: 801h/2049d Inode: 131490     Links: 1
Access: (0644/-rw-r--r--)  Uid: ( 1000/  jeremy)   Gid: ( 1000/  jeremy)
Access: 2019-10-13 15:26:39.567370013 +0200
Modify: 2019-10-13 15:26:39.567370013 +0200
Change: 2019-10-13 15:26:39.567370013 +0200
Birth: -
```

Then we rerun the command used before.

The `stat` command shows us for the screenshot :

```
stat /home/jeremy/Pictures/Screenshot\ from\ 2019-09-25\ 22-01-01.png
  File: /home/jeremy/Pictures/Screenshot from 2019-09-25 22-01-01.png
  Size: 445774      Blocks: 872      IO Block: 4096   regular file
Device: 801h/2049d Inode: 131311     Links: 1
Access: (0664/-rw-rw-r-- )  Uid: ( 1000/  jeremy)   Gid: ( 1000/  jeremy)
Access: 2019-10-13 13:36:16.777078803 +0200
Modify: 2019-09-25 22:01:01.653816263 +0200
Change: 2019-09-25 22:01:01.653816263 +0200
Birth: -
```

And for the file :

```
stat /home/jeremy/a
  File: /home/jeremy/a
  Size: 0      Blocks: 0      IO Block: 4096   regular empty file
Device: 801h/2049d Inode: 131490     Links: 1
Access: (0644/-rw-r--r-- )  Uid: ( 1000/  jeremy)   Gid: ( 1000/  jeremy)
Access: 2019-10-13 15:30:22.943370013 +0200
Modify: 2019-10-13 15:26:39.567370013 +0200
Change: 2019-10-13 15:26:39.567370013 +0200
Birth: -
```

We can see that the date accessed on the new file has changed, as for the old file, access has not changed.

6. **What happens to the files in the incremental backup that were hardlinked to the files of the full backup?**

## Task 2 : Set up SSH for remote login

1. In your personal `.ssh` directory download the `key.sec` file below this document. Be sure to remove all permissions for `group` and `others` from this file.

All permissions were removed using the `chmod` command.

2. Test logging into your account on the remote machine using SSH. Log out again.

```
ssh ait.lan.iict.ch -l jeremy_zerbib -i .ssh/key.sec
```

In order to log out, we used `exit` in the terminal.

3. On your local machine configure an SSH shortcut to the account on the remote machine. Create the file `~/.ssh/config` if does not yet exist and add the following lines to it:

```
# Cloud virtual machine for AIT lab
Host cloudvm
  Hostname ait.lan.iict.ch
  IdentityFile ~/.ssh/key.sec
  User jeremy_zerbib
```

In order to do this part, we copied the config above using :

```
gedit .ssh/config
```

Replace the username after `User` by your account name.

Test this shortcut by typing `ssh cloudvm`. You should see the command line prompt of the remote machine.

```
ssh cloudvm
Welcome to Ubuntu 16.04.6 LTS (GNU/Linux 4.4.0-165-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

5 packages can be updated.
0 updates are security updates.

New release '18.04.2 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Sun Oct 13 15:50:07 2019 from 10.192.18.127
```

## Task 3 : Remote Sync

1. **Create a backup directory on the remote machine as described in Task 1 so that your user can read/write.**

After login into the `ssh` session, we created a folder

```
mkdir /jeremy_backup
```

2. **Repeat the full backup and the incremental backup of task 1, but with the backup going to the remote machine over SSH. In the `rsync` command you need to prefix the destination parameter with `cloudvm:` to tell `rsync` to use SSH to transfer the data to the remote machine.**

```
rsync -av /home/ cloudvm:jeremy_backup/2017-09-25-093533 --delete --link-dest=2017-09-25-093533/
```

3. **Optional: Using a network monitoring tool on your local Linux machine like `bmon` observe how much network traffic `rsync` is causing.**

The bandwidth used by `ssh` is monitored by `bmon`.



